

2024 NTU Virtual Machine HW4 Writeup

StudentID : R12922054

Idea & Concepts

In this assignment, the objective is to execute a program named *hijack* within the virtual machine (VM). *hijack* attempts to overwrite its own code located in the virtual address range 0x4005e4 to 0x40060b. However, the memory page containing this code is marked as read-only, causing any write attempt to trigger a page fault. This assignment aims to modify the system so that the virtual address range 0x4005e4 to 0x40060b becomes writable, thereby allowing the *hijack* program to overwrite its code successfully.

The initial approach considered was to modify the TLB entries corresponding to the virtual address range 0x4005e4 to 0x40060b from being marked as read-only to writable. Based on insights from the resource [1], this approach seemed viable. Consequently, the focus shifted to identifying the function responsible for setting TLB attributes. The following sections detail the implementation process.

Implementation

1. First, find this function `arm_cpu_tlb_fill` in `/target/arm/tlb_helper.c`.
2. Second, in the function, `arm_cpu_tlb_fill`, it was observed that the variable, `prot`, determines the permissions for the memory address. Upon completing the execution of the function shown below, the value of `prot` would be set.

```
ret = get_phys_addr(&cpu->env, address, access_type,
                    core_to_arm_mmu_idx(&cpu->env, mmu_idx),
                    &phys_addr, &attrs, &prot, &page_size,
                    &fi, &cacheattrs);
```

3. Third, in the `/target/arm/helper.c` file, within the `get_phys_addr` function, we can see that the `get_phys_addr_lpae` function is executed, shown below. This function determines the `s2_prot` variable, which is then subjected to an AND operation with the `prot` variable. Therefore, we need to delve into the `get_phys_addr_lpae` function to perform operations on the `s2_prot` variable.

```
/* S1 is done. Now do S2 translation. */
ret = get_phys_addr_lpae(env, ipa, access_type, s2_mmu_idx, is_el0,
                          phys_ptr, attrs, &s2_prot,
                          page_size, fi, &cacheattrs2);
```

4. In the `/target/arm/helper.c` file, within the `get_phys_addr_lpae` function, the following code can be observed:

```
if (!(*prot & (1 << access_type))) {
    goto do_fault;
}
```

If we perform a `memcpy` operation, then `access_type = MMU_DATA_STORE` (where `access_type` has three values: `MMU_DATA_LOAD = 0`, `MMU_DATA_STORE = 1`, and `MMU_INST_FETCH = 2`). However, in the virtual address range 0x4005e4 to 0x40060b, the memory is read-only, meaning the value of `*prot` is `PAGE_READ`. Therefore, we need to add the following code before the code snippet shown above:

```
if ((address >= 0x4005e4 && address <= 0x40060b)){
    *prot |= PAGE_WRITE;
}
```

5. This way, we can successfully avoid entering the `if (!(*prot & (1 << access_type)))` block.

References

[1] "QEMU-TLB," [Online]. Available: https://airbus-seclab.github.io/qemu_blog/tcg_p3.html